
Timestamping algorithm

Topic 6, Lesson 6

An alternative to a locking protocol for concurrency control

Timestamping Algorithm

Concurrency control solution: alternative to a locking protocol, no locks so no deadlocks

Borrow an idea from the deadlock prevention algorithms:
place an order on the transactions, use the time order of the transactions to decide what operations are allowed in the database.

Conflict is resolved by rolling back and restarting transaction.
Older transactions, transactions with **smaller** timestamps, get less of a priority in the event of conflict.

Timestamp for transactions

Each transaction is assigned a unique timestamp

A timestamp is a unique identifier created by DBMS that indicates relative starting time of a transaction.

Can be generated by using system clock at time transaction started, or by incrementing a logical counter every time a new transaction starts.

Transaction's timestamp will indicate the order that the conflicting operations in the transactions are executed.

Timestamps for data objects

Also timestamps for data items:

read-timestamp - timestamp of last (most recent)
transaction to read item

write-timestamp - timestamp of last (most recent)
transaction to write item

When performing I/O operations, the read/write proceeds only if **last update on that data item** was carried out by an older transaction.

Otherwise, transaction requesting read/write is restarted and given a new timestamp, since it is considered stale

Timestamping algorithm

Two versions:

Basic timestamp ordering (BTO) algorithm – ensures that the timestamp for the transactions determines the order that the conflicting operations are completed (conflict serializable)

Thomas write rule (TW) algorithm – provides an optimization for dealing with stale updates of a data object. TW identifies when stale updates can be safely ignored and not completed

Neither version guarantees a recoverable schedule

Timestamping Algorithm: Read(x)

Consider a transaction T with timestamp $ts(T)$:

$ts(T) < write_timestamp(x)$

x already updated by younger (later) transaction.

Transaction must be aborted and restarted with a new timestamp.

Otherwise $ts(T) > write_timestamp(x)$

operation is accepted and executed.

Update the $read_timestamp(x)$

$read_timestamp = \max((ts(T), read_timestamp(x)))$

Timestamping Algorithm: write(x) BTO

ts(T) < read_timestamp(x)

x already read by a younger transaction.

Roll back transaction and restart it using a later timestamp.

ts(T) < write_timestamp(x)

x already written by younger transaction. T's version of x is obsolete.

Roll back transaction and restart it using a later timestamp.

Otherwise, operation is accepted and executed. Set the
write_stamp = ts(T)

Timestamping Algorithm: write(x) (TW)

ts(T) < read_timestamp(x)

x already read by a younger transaction.

Roll back transaction and restart it using a later timestamp.

ts(T) < write_timestamp(x)

x already written by younger transaction. T's version of x is obsolete.

Write can safely be ignored - **ignore obsolete write Thomas' write** rule.

Otherwise, operation is accepted and executed. Set the **write_stamp = ts(T)**

TW rule: Simple Example

Time	Transaction 1	Transaction 2
t ₁	Start Transaction	
t ₂		Start Transaction
t ₃		Read(A)
t ₄	Read(B)	
t ₅	Ignored	Write(C)
t ₆	Write(C)	
t ₇	Commit	
t ₈		Commit

Time	Transaction 1	Transaction 2
t ₁	Start Transaction	
t ₂		Start Transaction
t ₃		Read(A)
t ₄	Read(B)	
t ₅		Write(C)
t ₆		
t ₇	Commit	
t ₈		Commit

Timestamping algorithm: Example

Time	Op	T ₁₉	T ₂₀	T ₂₁
t ₁		begin_transaction		
t ₂	read(bal _x)	read(bal _x)		
t ₃	bal _x = bal _x + 10	bal _x = bal _x + 10		
t ₄	write(bal _x)	write(bal _x)	begin_transaction	
t ₅	read(bal _y)		read(bal _y)	
t ₆	bal _y = bal _y + 20		bal _y = bal _y + 20	begin_transaction
t ₇	read(bal _y)			read(bal _y)
t ₈	write(bal _y)		write(bal _y) [†]	
t ₉	bal _y = bal _y + 30			bal _y = bal _y + 30
t ₁₀	write(bal _y)			write(bal _y)
t ₁₁	bal _z = 100			bal _z = 100
t ₁₂	write(bal _z)		T ₂₂	write(bal _z)
t ₁₃	bal _z = 50	bal _z = 50		commit
t ₁₄	write(bal _z)	write(bal _z) [‡]	begin_transaction	
t ₁₅	read(bal _y)	commit	read(bal _y)	
t ₁₆	bal _y = bal _y + 20		bal _y = bal _y + 20	
t ₁₇	write(bal _y)		write(bal _y)	
t ₁₈			commit	

[†] At time t₈, the write by transaction T₂₀ violates the first timestamping write rule described previously and therefore is aborted and restarted at time t₁₄.

[‡] At time t₁₄, the write by transaction T₁₉ can safely be ignored using the ignore obsolete write rule, as it would have been overwritten by the write of transaction T₂₁ at time t₁₂.

Disadvantages to timestamping algorithm

If a transaction is aborted, it is restarted with a new timestamp. This can result in a cyclic restart where a transaction can **repeatedly restart and abort without ever completing**, penalizes long-running transactions.

Another disadvantage is that it has storage overhead for maintaining timestamps (two timestamps must be kept for every data object).

Also requires more write operations to the database since we are tracking and updating each access (read or write) to the database objects.

Summary

- Timestamp is another method for supporting concurrency
- It is a simpler solution to locking since it only requires tracking read and write access to data objects and placing an ordering on the transactions
- It can suffer from starvation of long-running transactions
- It requires more updates to the database, since it does an update of the read timestamp when transaction's read the data.